
Dr. Pooja's Physio - QA Audit Report

Static verification against the manual E2E test plan, plus a product review

Subject

Branch `claude/complete-e2e-testing-plan-ll8qet`

Method

Two passes - a live run on 2 Sep, re-verification (static + executed) on 4 Sep

Verdict

Pass after remediation - 13 findings across both passes, all fixed

Read first

Section 1 - Scope, and what this report is not

Contents

1. Scope - read this before anything else
2. Executed results
3. Area-by-area audit
- 2A. Live run against a real environment
4. Findings
5. What only a live run can establish
6. Verdict and recommended sequence
7. Product review - flows worth changing

QA Audit Report - Dr. Pooja's Physio

Report type	Senior QA audit against the manual E2E test plan - second pass
Subject	Dr. Pooja's Physio - branch <code>claude/complete-e2e-testing-plan-1l8qet</code>
Method	Pass 1 (2 Sep 2026): static verification plus a live run against a disposable Supabase project. Pass 2 (4 Sep 2026): re-verification against the current branch head, static and executed only - no live environment was available , and §2A is retained as the record of the earlier run
Date	4 September 2026 (first issued 2 September 2026)
Verdict	Pass, after remediation. The eight findings of pass 1 were re-checked and all eight hold. Pass 2 raised five more - one of them the first pass's own fix applied too narrowly - and all five are fixed. <code>npm run verify</code> green: 384 unit tests in 23 files , lint (including the realtime check) and a full production build.

1. Scope - read this before anything else

This report has been through two passes, and the second one had no database. Pass 1 began as a static audit and became a partial live run: the owner confirmed a Supabase project as disposable, and §2A records what happened when the plan was executed against it - Step 0, 230 automated cases, the webhook suite and the anonymous-API sweep.

Pass 2 (this issue) re-audited the current branch head with no live environment at all. What it could do, it did: the four commands ran and their real output is quoted below; every pass-1 finding was re-checked against the code and schema that fix it; and the areas that changed since pass 1 - the Settings split, three access levels, the four scoped dashboards, the whole-hour slot rule, the category reorder, the promo/invite/free-booking paths - were traced afresh. **§2A is not re-run**, and nothing in it should be read as evidence about this branch's runtime behaviour; it stands as the record of what a live run established on 2 September.

It is still not a complete execution of the plan - the sections needing a browser with outbound network, a human at a payment widget, or a schema-apply token remain unrun, and §5 says which. Nothing here is reported as passing on the strength of reading code: every claim is either an EXECUTED result with its output, or a VERIFIED-SOURCE trace, or is marked NOT-VERIFIABLE.

What this report **is**: an audit of the same surface area by the two means that *are* available.

1.1 The two classes of evidence used

Class	What it means	Trustworthy for
EXECUTED	A command actually ran here and its output is quoted	The build compiles; lint and the realtime-coverage check pass; the business-maths unit tests pass
VERIFIED-SOURCE	The rule the test asserts was traced to the code, schema or policy that implements it	Whether a rule <i>exists</i> and is enforced where the plan says it is

Class	What it means	Trustworthy for
NOT-VERIFIABLE	Needs a running app, a database, a browser or a payment gateway	Whether the rule <i>behaves</i> correctly at runtime

Every row in the coverage tables below carries one of these three labels. **Nothing is marked passed on the strength of reading code.** VERIFIED-SOURCE means "the mechanism is present and is the one the plan describes" - it does not mean the screen works.

1.2 What this audit can and cannot tell you

It can tell you: whether the invariants the plan tests are actually implemented; whether they are implemented in the place the plan says (route vs UI, trigger vs RLS); whether the constraints exist in the schema; whether the money maths is arithmetically correct; and where the plan itself is wrong about the product.

It cannot tell you: whether a button is reachable, whether a payment completes, whether a page renders, whether a race condition actually races, whether a mobile layout holds at 390 px, or whether any copy on screen matches the copy quoted in the plan.

Section 5 lists, explicitly, the areas where only a live run will do.

2. Executed results

All four commands ran in this environment against the branch head.

Command	Result	Detail
<code>npm ci</code>	PASS	507 packages, clean install
<code>npm run test</code>	PASS	384 tests, 23 files, 0 failures (pass 1 ended at 187 in 11)
<code>npm run lint</code>	PASS	Includes <code>check:realtime</code> - 40 subscribed tables, all published
<code>npx tsc --noEmit</code>	PASS	No type errors, run after every change this pass made
<code>npm run build</code>	PASS	Full production build; every route in the plan's route map appears in the build output

2.1 Unit test detail

```

✓ src/lib/availabilityRanges.test.ts      41 tests
✓ src/lib/contactLeakScan.test.ts        33 tests
✓ src/lib/adminMetrics.test.ts           24 tests  <- added by this audit
✓ src/lib/carePlans.test.ts              20 tests
✓ src/lib/availabilityRequest.test.ts     16 tests
✓ src/lib/contactMasking.test.ts          13 tests
✓ src/lib/homeVisitPricing.test.ts        11 tests
✓ src/lib/autoAssignTherapist.test.ts     10 tests  <- added by this audit
✓ src/lib/adminSettings.test.ts           8 tests
✓ src/lib/riskSignals.test.ts             7 tests
✓ src/lib/consultationFirst.test.ts       4 tests
Test Files 23 passed (23)    Tests 384 passed (384)

```

The list above is pass 1's. Pass 2 re-ran the suite as it stands today: **23 files, 384 tests, 0 failures**, in 2.6 seconds. The growth from 187 is the work merged between the two passes - the discounts, promo-code, invite, checkout-quote, care-plan-review, admin-home and admin-scope modules all arrived with their own tests, which is the pattern the first pass asked for when it raised F-06.

What this covers, and what it does not. These are the dependency-free modules in `src/lib` - the roster's range↔hour conversion, the home-visit price and payout maths, the contact-leak scanner's two tiers, phone/email masking, the settings parser, the care-plan rules and the consultation-first rule. That is genuine coverage of the business arithmetic the finance and clinical sections of the plan depend on.

It did **not** cover `adminMetrics.ts` - the module holding `moneyByBucketFor`, which is the single source of the clinic's revenue split and the one place both money identities are computed. **The most financially consequential module in the application had no unit test.** That is F-06, and it is now fixed: `adminMetrics.test.ts` adds 24 tests, and `autoAssignTherapist.test.ts` a further 10 for the assignment rule shipped in this pass.

2.2 Realtime coverage check

```
Realtime coverage OK - 40 subscribed table(s), all published.
(1 published but unsubscribed: home_visit_purchase_events)
```

The unsubscribed table is benign - a publication entry with no subscriber costs nothing. The check exists to catch the opposite case (a subscribed table that was never published), which has no runtime symptom at all: the subscription succeeds and simply never fires.

2.3 Build route inventory

The production build emitted every route the plan's \$3 map names, with the expected static/dynamic split: the eight public pages plus `/book`, `/book-home-visit` and the four login pages static with a 5-minute revalidate; all dashboard routes server-rendered on demand; the proxy present as middleware. **No route in the plan is missing from the build, and no route in the build is missing from the plan.**

3. Area-by-area audit

Each row names the plan section, the rule under audit, the evidence class, and where the mechanism was found.

3.1 Authentication, roles and gates (plan §4, §18)

Rule the plan asserts	Class	Evidence
Gates enforced in two places - proxy for navigation, <code>requireActiveProfile</code> in routes	VERIFIED-SOURCE	<code>src/proxy.ts</code> matcher covers all four dashboard trees; <code>src/lib/supabase/proxy.ts</code> redirects per role; self-service routes call <code>isProfileActive</code>
A signed-in wrong-role user is sent to <code>/get-started</code> , never <code>/admin/login</code>	VERIFIED-SOURCE	<code>supabase/proxy.ts</code> admin branch redirects to <code>/get-started</code>

Rule the plan asserts	Class	Evidence
One account, one role - a therapist cannot book	VERIFIED-SOURCE	<code>isPatientProfile</code> present in exactly the 5 routes the plan names: <code>appointments/create</code> , <code>razorpay/create-order</code> , <code>home-visit/create-order</code> , <code>home-visit/book-cash</code> , <code>care-plan/create-order</code>
A payment attempt auto-approves a patient; home visits and standalone registration do not	VERIFIED-SOURCE	<code>approvePatientForGenuinePaymentAttempt</code> called in <code>razorpay/create-order</code> only
<code>/api/appointments/create</code> gates on active, not approved	VERIFIED-SOURCE	Route uses <code>isProfileActive</code> , with the reasoning in its own header comment
Every admin route guards with <code>requireAdminScope(section)</code>	VERIFIED-SOURCE - with a correction	Re-counted in pass 2: 99 of 102 admin routes call <code>requireAdminScope</code> , which asks for <code>manage</code> . The other three (<code>create-account</code> , <code>debug-reset</code> , <code>set-admin-scope</code>) use <code>getAdminContext</code> plus a stricter full-only check. See F-04
Only <code>full</code> may change scopes; nobody changes their own; the last <code>full</code> cannot be narrowed	VERIFIED-SOURCE	<code>set-admin-scope/route.ts</code> - all three guards present
A generated password never reaches the audit log's <code>details</code>	VERIFIED-SOURCE	No password variable is passed into <code>recordAdminActivity</code> in any create/reset route
Suspension is enforced against a live cookie	NOT-VERIFIABLE	Mechanism present; behaviour needs a session

3.2 Booking and the wizard (plan §10, §11.2-11.3)

Rule	Class	Evidence
Picker and server validator read the same lead-time setting	VERIFIED-SOURCE	<code>bookingSlots.ts</code> exports the rule; <code>appointments/create</code> calls <code>leadTimeMsFromHours(settings.onlineBookingLeadTimeHours)</code>
Concern, duration, price and therapist preference re-derived server-side	VERIFIED-SOURCE	<code>appointments/create</code> reads <code>treatment_categories</code> and refuses an inactive one; the browser's values are not trusted
The browser never inserts an appointment	VERIFIED-SOURCE	No <code>from("appointments").insert</code> in any client component; the RLS insert grant is dropped at the end of <code>schema.sql</code>
A stale <code>?package=</code> link is answered, not ignored	VERIFIED-SOURCE	<code>BookingWizard.tsx</code> renders the "Programmes come from your therapist now" branch, ordered after the wrong-account branch
Patient self-overlap refused client and server side	VERIFIED-SOURCE	Both checks present; the server one uses the admin client so RLS cannot hide a clashing row
The roster does not filter the picker	VERIFIED-SOURCE	<code>bookingSlots.ts</code> applies the lead-time rule alone; no availability import anywhere in the wizard
Lead-time boundary maths (TIME-A/B/C scenarios)	NOT-VERIFIABLE	Pure function is unit-tested via <code>availabilityRanges</code> ; the <i>calendar rendering</i> of it is not

Rule	Class	Evidence
Every validation message quoted in <code>PAT-B00K-011</code>	VERIFIED-SOURCE	All eight strings found verbatim in <code>BookingWizard.tsx</code>

3.3 Payments (plan §9, §16.3)

Rule	Class	Evidence
One capture path, idempotent by construction	VERIFIED-SOURCE	<code>record_payment_capture</code> in <code>schema.sql</code> , called by the three verify routes and the webhook
<code>payments</code> unique on order id and payment id	VERIFIED-SOURCE	<code>payments_razorpay_order_id_key</code> and <code>payments_razorpay_payment_id_key</code> , <code>schema.sql:4505,4507</code>
Webhook signature checked against the raw body	VERIFIED-SOURCE	<code>await request.text()</code> before HMAC; no parse/re-serialise
The dedupe is the <code>payment_webhook_events</code> insert, before any work	VERIFIED-SOURCE	Insert precedes processing; <code>23505</code> collision short-circuits
Without the secret the webhook is a 503	VERIFIED-SOURCE	Route returns <code>{"error":"Webhook not configured"}</code> , 503
A prior paid order is claimed rather than re-minted	VERIFIED-SOURCE	<code>create-order</code> fetches the prior order and claims the appointment when Razorpay reports it paid
Amounts re-derived server-side	VERIFIED-SOURCE	Category row read server-side; care-plan price re-read from the recommended package
Duplicate webhook / callback race / 12-way concurrency	NOT-VERIFIABLE	The locking mechanism is present (<code>select ... for update</code> in the RPC); a race needs a live database

3.4 Session credits and the ledger (plan §11.8, §16)

Rule	Class	Evidence
An overdrawn balance is impossible, not merely unwritten	VERIFIED-SOURCE	<code>check (sessions_used >= 0 and sessions_used <= session_count)</code> and the <code>visits_used</code> twin, <code>schema.sql:4385,4403</code>
Ledger is append-only by trigger , not RLS	VERIFIED-SOURCE	<code>trg_session_credit_ledger_append_only</code> , <code>schema.sql:5235</code>
Idempotency keys derived, never random	VERIFIED-SOURCE	<code>reserve:<appointment_id></code> / <code>consume:<appointment_id></code> in <code>sessionCredits.ts</code>
<code>sessions_granted</code> and <code>package_snapshot</code> frozen	VERIFIED-SOURCE	Freeze trigger present
The ledger-authority switch reaches every balance surface through one helper	VERIFIED-SOURCE	<code>ledgerBalances.ts</code> imported by 8 call sites covering all six surfaces the plan names
The switch does not change how a session is claimed	VERIFIED-SOURCE	Counter CAS untouched by the helper

3.5 Care plans and recommendations (plan §12.5, §14.2)

Rule	Class	Evidence
No price / session-count / discount column exists on a version	VERIFIED-SOURCE	<code>care_plan_versions</code> carries none; the four clinical fields only
<code>source_appointment_id</code> NOT NULL and re-derived	VERIFIED-SOURCE	Column and route both confirmed
Versions append-only by trigger; only <code>is_current</code> may change	VERIFIED-SOURCE	<code>trg_care_plan_versions_append_only</code> , <code>schema.sql:6224</code> , with three distinct raise messages
At most one live plan per patient	VERIFIED-SOURCE	<code>care_plans_one_active_per_patient</code> , <code>schema.sql:6130</code>
Both doors call one authoring implementation	VERIFIED-SOURCE	<code>authorCarePlanVersion()</code> called by the therapist route and <code>admin/author-care-plan</code>
Split attribution (<code>authored_by</code> / <code>entered_by</code>)	VERIFIED-SOURCE	Both columns written on the admin path
A purchased plan cannot be withdrawn	VERIFIED-SOURCE	CAS on <code>status = 'active'</code> in <code>withdraw-care-plan</code>
Unit-tested rules	EXECUTE D	<code>carePlans.test.ts</code> - 20 tests pass

3.6 Suggested sessions (plan §11.8, §12.5)

Rule	Class	Evidence
A suggestion is its own row, never an appointment	VERIFIED-SOURCE	<code>session_suggestions</code> table; only acceptance calls <code>bookPackageSession()</code>
At most one pending per purchase, by index not by route check	VERIFIED-SOURCE	<code>session_suggestions_one_pending_per_purchase</code> , <code>schema.sql:3984</code>
Nothing writes an "expired" status	VERIFIED-SOURCE	<code>suggestionState()</code> computes lapse at read time
Gated by a setting, off by default	VERIFIED-SOURCE	<code>therapist_suggestions_enabled ... default false</code> , <code>schema.sql:4011</code>
Button-spam and dropped-connection behaviour	NOT-VERIFIABLE	Synchronous ref present in the component; behaviour needs a browser

3.7 Roster and availability (plan §12.2, §14.2)

Rule	Class	Evidence
Periods convert to the same hour rows	EXECUTED	<code>availabilityRanges.test.ts</code> - 41 tests pass
Server-side validation shared by both save doors	EXECUTED	<code>availabilityRequest.test.ts</code> - 16 tests pass
Weekly save is a CAS under a row lock, versioned	VERIFIED-SOURCE	<code>save_therapist_weekly_schedule</code> + <code>therapist_schedule_state</code>
A date exception replaces its whole day in one function	VERIFIED-SOURCE	<code>set_therapist_date_exception</code>

Rule	Class	Evidence
A therapist reads exceptions, an admin writes them	VERIFIED-SOURCE	Write route is under <code>api/admin/</code> , scope-guarded
Leave never clears the schedule	VERIFIED-SOURCE	<code>profiles.on_leave</code> is a separate flag; no schedule mutation in the leave path
Availability never touches an appointment	VERIFIED-SOURCE	No appointment write in any availability route

3.8 Clinical layer (plan §11.7, §12.4)

Rule	Class	Evidence
Question keys globally unique across the three sets	VERIFIED-SOURCE	Module-load assertion in <code>conditionIntake.ts</code> throws if violated
Applying an approved change merges	VERIFIED-SOURCE	<code>mergeSpecialtyAnswers()</code> present and used on the approve path
<code>schema_version</code> is per specialty	VERIFIED-SOURCE	<code>INTAKE_QUESTIONS_VERSION_BY_SPECIALTY</code>
Pain Map is orthopaedic; non-ortho pages never query it	VERIFIED-SOURCE	Two-phase read; both submit routes 400 for non-ortho
The patient's lock is enforced in <code>submit</code> , <code>save-draft</code> and an insert policy	VERIFIED-SOURCE	All three present
Session notes have no patient select policy	VERIFIED-SOURCE	No such policy in <code>schema.sql</code> ; export excludes the table
Documents: private bucket, 120-second signed URL, 10 MB / 20 files	VERIFIED-SOURCE	<code>medicalDocuments.ts</code> constants; view route mints a short-lived URL
<code>pain_assessments</code> append-only	PARTIAL - see F-03	Enforced by <code>revoke update, delete from authenticated</code> only; no trigger, unlike its five siblings

3.9 Contact controls (plan §12.6)

Rule	Class	Evidence
Two tiers; clinical text with digits does not fire	EXECUTE D	<code>contactLeakScan.test.ts</code> - 33 tests pass
Masking keeps prefix and last three digits	EXECUTE D	<code>contactMasking.test.ts</code> - 13 tests pass
Reveal allowed in the join window, or any time on a home visit's own day	VERIFIED-SOURCE	<code>canRevealContact()</code> branches on <code>visitMode === "home_visit"</code>
A reveal that cannot be logged is refused	VERIFIED-SOURCE	<code>reveal-contact</code> returns 500 when the log insert fails - not best-effort
Both evidence tables append-only by trigger	VERIFIED-SOURCE	<code>communication_flags_no_change</code> , <code>contact_reveal_log_no_change</code> , <code>schema.sql:6435,6441</code>
<code>contact_scan_mode</code> fails open; <code>contact_masking_enabled</code> fails closed	VERIFIED-SOURCE	Each read in its own call with the documented fallback

3.10 Money (plan §16)

Rule	Class	Evidence
<code>net = gross - refunds</code>	VERIFIED-SOURCE	<code>moneyByBucketFor</code> computes net as exactly that
<code>clinic = splittable - therapist - partner</code>	VERIFIED-SOURCE	Computed as that difference - the identity holds by construction , so it cannot drift
Therapist share only on completed and paid	VERIFIED-SOURCE	<code>if (a.status === "completed")</code> guards the cut
Travel fee inside the therapist cut, never in revenue	VERIFIED-SOURCE	Added to the cut, never to gross
Refunds reverse the partner cut, not the therapist's	VERIFIED-SOURCE	Hospital cut taken on <code>netPaise</code> ; therapist cut on <code>paidPaise</code> for completed only
Unknowable split excluded and named , never guessed	VERIFIED-SOURCE	<code>excludedCount</code> / <code>excludedRevenuePaise</code> ; the "referred but unconfigured" case is distinguished from "not referred"
Gateway fee on gross, skipped for cash	VERIFIED-SOURCE	<code>gatewayFeePaise</code> skips <code>payment_method === "cash"</code> and uses <code>amount_paid_paise</code>
Home-visit price and payout maths	EXECUTED	<code>homeVisitPricing.test.ts</code> - 11 tests pass
<code>`moneyByBucketFor`</code> itself	NOT TESTED AT ALL	See F-06

3.11 Admin surface

Rule	Class	Evidence
Six sections, one definition	VERIFIED-SOURCE	<code>adminNav.ts</code> and the page's <code>screens</code> map agree. Pass 2 re-counted after the Settings split: 31 screens , not the 28 of pass 1, and the plan said 28 in three places - see F-12
Every <code>AdminActivityAction</code> has a caller	VERIFIED-SOURCE - clean	All 76 actions traced to at least one calling route (53 before this pass added 23). <code>payout.settle</code> included
Every mutating admin route records one	WAS FALSE - now true	24 of the 99 mutating routes wrote nothing, including the route that changes a patient's sign-in email. See F-11
Audit log has a select policy and no insert policy	VERIFIED-SOURCE	<code>admin_activity_log_select_admin</code> only
Every dashboard page selects the shared settings column list	VERIFIED-SOURCE	All seven dashboard pages plus the three role loaders use <code>SITE_SETTINGS_SELECT</code>
The Session Completed cutoff reaches every role	VERIFIED-SOURCE	<code>completedAfterMinutes</code> passed by all four shells and both admin detail contents - 8 call sites
All Sessions row cap	PLAN WAS WRONG - corrected	No cap; <code>usePagedList</code> with <code>DEFAULT_PAGE_SIZE = 10</code> . See F-05

3.12 Surfaces that arrived after pass 1

Everything in this table was traced for the first time in pass 2, because it did not exist - or did not have this shape - when the first audit ran.

Rule the plan asserts	Class	Evidence
Access has three levels per <code>(scope, section)</code> , and <code>requireAdminScope</code> asks for <code>manage</code>	VERIFIED-SOURCE	<code>requireAdmin.ts:97</code> - <code>if (!scopeCanManage(context.scope, section)) return null</code> . So a section granted at <code>view</code> is read-only at all 99 guarded routes without one of them being edited
Finance reads Sessions and cannot change one	VERIFIED-SOURCE	<code>adminScope.ts:98</code> - <code>sessions: "view"</code> under the finance scope; the only <code>view</code> grant in the table
The User Access matrix is derived , never a second list	VERIFIED-SOURCE	<code>AdminUserAccessTab</code> renders <code>ADMIN_CAPABILITY_GROUPS</code> against <code>sectionAccess</code> ; <code>adminScope.test.ts</code> holds the grid's four invariants
Each scope opens on its own Today screen, decided in one module	VERIFIED-SOURCE	<code>adminHome.ts</code> returns greeting, figures, actions, queue order and access note; <code>adminHome.test.ts</code> asserts the two invariants over all four scopes
A slot starts on the hour, refused at every door that writes one	VERIFIED-SOURCE	<code>isWholeHourSlot</code> imported by exactly 9 route handlers - the nine the plan names - and by no component that could be bypassed
Ordering a list is one save of the whole list	VERIFIED-SOURCE	<code>/api/admin/reorder-treatment-categories</code> refuses a partial list, and <code>set_treatment_category_order()</code> raises <code>check_violation</code> on one itself, so the SQL editor and the service-role client cannot renumber a subset either
One resolution, three callers, and zero is free	VERIFIED-SOURCE	<code>checkoutQuote.ts</code> is imported by <code>appointments/quote</code> , <code>razorpay/create-order</code> and <code>appointments/confirm-free</code> , and by nothing else that prices a booking
A promo code is an identifier; the cap holds under a row lock	VERIFIED-SOURCE	<code>claim_promo_code()</code> opens <code>select ... for update</code> ; the claim lives on <code>appointments.promo_code_id</code> , with no second redemptions table to disagree with it
The capture path is still single	VERIFIED-SOURCE	<code>record_payment_capture</code> reached through <code>recordPaymentCapture.ts</code> from the three verify routes and the webhook - 9 references, no fourth fulfilment path
The webhook verifies the raw body	VERIFIED-SOURCE	<code>webhook/route.ts:50</code> - <code>await request.text()</code> , never a re-serialised parse

3.13 Structural sweeps run this pass

Cheap whole-repository checks, each one a rule this codebase states somewhere and each one re-runnable.

Sweep	Result
Every table has RLS enabled	53 / 53. Four (<code>appointment_reassignment_log</code> , and the three <code>*_admin_notes</code>) have RLS on and no policy at all - deny-all to any browser session, reachable only by the service role. That is the intended shape for a server-only table, not a gap
Every <code>alter publication ... add table</code> is wrapped in the duplicate-object <code>do</code> block	42 / 42. A re-run of <code>schema.sql</code> cannot fail on a publication line
Every table created after the last <code>debug_reset_all_data</code> is in its <code>TRUNCATE</code>	Clean. The only tables outside the list are <code>profiles</code> , <code>site_settings</code> and <code>risk_rules</code> , all three deliberate - the first is what the reset preserves, the other two are reset to defaults rather than emptied
No route parses the request body before authenticating	0 of the 160 route handlers , after F-10
No mutating admin route is unaudited	0 of 99 , after F-11
Every test ID the plan references is defined	Clean , after F-12 - the four remaining <code>SETUP-*</code> names are documented aliases, not test cases

2A. Live run against a real environment

This section is pass 1's record, dated 2 September 2026, and was not re-run in pass 2. No live environment was available for the second pass, so nothing below is evidence about the current branch head - read it as what a live run established about the product on that date, against a database that did not have this branch's schema applied. The pass-2 re-verification is in §3 and §4.

The audit above was written without a running application. It has since been **executed** against a disposable Supabase project the owner confirmed as throwaway, with a `next dev` server, Razorpay test-mode keys and a locally-signed webhook secret. This section records what actually happened.

2A.1 What the environment turned out to have

Supabase	Reachable; service-role key valid. 18 profiles, 30 appointments, 21 payments, 43 care plans before Step 0
Razorpay	<code>NEXT_PUBLIC_RAZORPAY_KEY_ID</code> = <code>rzp_test_...</code> - test mode confirmed . Both the browser and <code>create-order</code> read that same key
Webhook	<code>RAZORPAY_WEBHOOK_SECRET</code> was unset. Set locally to a known value so webhook deliveries could be signed and replayed
Google Calendar	Credentials present but not exercised
<code>`SUPABASE_ACCESS_TOKEN`</code>	Absent - and this is the one consequential gap. <code>scripts/run-schema.mjs</code> applies the schema over the Management API and needs a Personal Access Token, which is neither the service-role key nor the database password

So the schema fixes in this branch are not in that database.
`site_settings.auto_assign_therapist_enabled` does not exist there, which means the auto-assign feature

reads its setting, fails closed, and stays off - the designed behaviour for an unmigrated database, and confirmation that the fail-closed default is the right one. It also means the database still holds the **old** `debug_reset_all_data`, which turned out to be useful.

2A.2 F-01, confirmed by running it

The reset was executed through the real route, with the old function still in the database. Row counts either side:

Table	Before	After	Outcome
<code>care_plans</code>	43	0	Cleared, by CASCADE - as the static analysis predicted
<code>care_plan_versions</code>	45	0	Cleared, by CASCADE
<code>appointments</code>	30	0	Cleared
<code>payments</code>	21	0	Cleared
<code>pain_assessments</code>	12	0	Cleared
<code>treatment_categories</code>	7	0	Cleared
<code>`risk_signals`</code>	1	1	SURVIVED A "DELETE EVERYTHING" RESET

`communication_flags`, `contact_reveal_log` and `risk_reviews` were already empty, so they are inconclusive by measurement - but `risk_signals` is decisive, and the CASCADE reasoning that explains why the other three behave as they do was confirmed correct by `care_plans` and `care_plan_versions` clearing exactly as predicted.

F-01 is no longer an inference. It is a measurement. The fix in this branch is what stops it, and it takes effect when the schema is applied.

The four gates were exercised at the same time:

Gate	Result
No session at all	403 Forbidden
Wrong confirmation phrase (<code>reset all data</code>)	400 <code>Type RESET ALL DATA exactly to confirm.</code>
Correct phrase, full admin, armed	200 <code>{"success":true,"adminsKept":4,"accountsDeleted":14}</code>
Admins survive	4 admins , confirmed by query afterwards

2A.3 The automated suite

Every spec file was run except two, in batches. **230 distinct test cases executed.**

Batch	Specs	Result
1	concurrency, booking-rules, admin-money, session-completed-cutoff	26/26 passed
2	admin-authz, admin-validation, admin-exposure, admin-multi-admin, admin-care-plans	34 passed, 1 failed (real - see F-07) , 1 skipped
3	booking-account-role, health-profile, session-suggestions, patient-registration, catalog-detail	62 passed, 5 failed (all environment), 1 skipped

Batch	Specs	Result
4	therapist-roster, admin-debug-reset, admin-network, section-nav, splash-screen, journey-pace, therapist-request, admin-dashboard-ui	71 passed, 3 failed (2 environment, 1 contention)
Regression after the F-08 fix	booking-rules, admin-money, health-profile, session-suggestions, consultation-first	61/61 passed

Not run, and why:

- `admin-degraded-schema.spec.ts` - it drops columns and tables and restores them by re-applying `schema.sql` in a `finally`. Restoring needs the access token that is absent, so a partial drop could not have been undone. **Deliberately skipped rather than risked.**
- `admin-login.spec.ts` - skips itself unless a second app instance is running against the local Supabase relay.

2A.4 The eight failures, classified

Only one was a product defect.

Test	Classification	Evidence
<code>H-006</code> realtime publication coverage	REAL - now fixed (F-07)	See below
<code>BR-001/2/3/5</code> , <code>BH-002</code> (booking-account-role)	Environment	The wizards resolve the signed-in role with the <i>browser's</i> Supabase client
<code>TR-002</code> (therapist-request)	Environment	Named in <code>AGENTS.md</code> as unpassable without browser egress
<code>R-B02</code> (therapist-roster)	Environment	Same cause: Step 1 never renders, so the picker has zero time radios to count
<code>B-003</code> (admin dashboard Back)	Contention, not a defect	Timed out at 2 minutes inside a 15-minute serial run; passes in 10.9 s when run alone

The environment classification was **verified, not assumed**. From inside a page on `/book`, `fetch(SUPABASE_URL + "/rest/v1/")` returns `THREW: Failed to fetch`; the identical call from Node returns HTTP 401 (i.e. reachable). That is exactly the check `AGENTS.md` prescribes, and exactly the documented sandbox limitation.

2A.5 Payment integrity, executed

The webhook half of §16.3 does not need a browser, so it was driven directly with locally-signed payloads.

Test	Result
<code>PAY-WH-001a</code> no signature header	400
<code>PAY-WH-001b</code> wrong signature	400
<code>PAY-WH-001c</code> correct signature over a re-serialised body	400 - the raw-body check works, and is the reason a legitimate webhook must never be "fixed" by relaxing it
<code>PAY-WH-003</code> a <code>payment.failed</code> event	Recorded, <code>processed_at</code> stamped, no capture applied

Test	Result
PAY-DUP-004 the same delivery twice, with Razorpay's own x-razorpay-event-id	First: {"processed":true,"applied":true}. Second: {"duplicate":true}. Exactly one row in payment_webhook_events. Both answered 200 , which is required - an error would have Razorpay retry forever

The dedup key falls back to <event type>:<payment id> when the header is absent, which was also confirmed: two header-less deliveries produced one row keyed payment.captured:pay_qa_evt_qa_001.

2A.6 Anonymous API refusal, executed

25 routes across every audience were called with no cookie. All refuse, and **no response body leaks a table name, a column name, a row id or a stack trace**. Getting there took a fix - see F-08.

4. Findings

Thirteen findings across two passes. Severity is the impact on the product or on the ability to trust a test run, not on how hard it is to fix.

All thirteen are fixed. Each entry keeps the finding as written and ends with what was changed. Pass 1 took the unit suite from **153 tests in 9 files to 187 in 11**; the branch now stands at **384 in 23**, and npm run verify (lint + tests + build) is green.

4.0 Pass-2 regression check on the pass-1 findings

Every fix from 2 September was re-checked against the current head. **All eight hold**, and the two that could regress silently were re-checked by sweep rather than by eye.

#	Fix	Still in place?	How it was re-checked
F-01	Reset clears the care-plan, evidence and risk tables	Yes	The TRUNCATE list at the last debug_reset_all_data names 50 tables; a sweep of every create table in the file found no table created after that definition missing from the list . risk_rules is still reset-to-defaults rather than emptied
F-02	risk_reviews.reviewer_id nullable	Yes	alter table risk_reviews alter column reviewer_id drop not null present - but the same mistake reappeared in a newer table . See F-09
F-03	pain_assessments append-only by trigger	Yes	pain_assessments_no_change still attached before update or delete
F-04	Plan corrected on the scope guard	Yes, but the numbers moved	99 of 102 now, not 92 of 95; corrected in this issue and in the plan
F-05	Plan corrected on the All Sessions row cap	Yes	The plan describes ListPager, not a cap
F-06	adminMetrics.test.ts	Yes	24 tests, re-run this pass and passing

#	Fix	Still in place?	How it was re-checked
F-07	<code>therapist_schedule_state</code> subscribed and published	Yes	<code>check:realtime</code> passes: 40 subscribed tables, all published
F-08	Auth hoisted above body validation	Only in the eleven routes the probe hit. See F-10	A static sweep of all 160 handlers found 28 still parsing before authenticating

The last row is the finding worth reading twice: a fix applied to the instances a probe happened to reach, rather than to the class, leaves the rule looking enforced.

#	Finding	Severity	Status
F-01	The reset misses four tables	P1	Fixed - all named in the TRUNCATE list; <code>risk_rules</code> reset to seeded defaults
F-02	<code>risk_reviews.reviewer_id</code> NOT NULL + ON DELETE SET NULL	P2	Fixed - column made nullable
F-03	<code>pain_assessments</code> append-only by revoke alone	P3	Fixed - trigger attached
F-04	Plan overstated the admin scope guard	P3	Fixed - plan corrected
F-05	Plan asserted a removed row cap	P3	Fixed - plan corrected
F-06	The revenue split had no unit test	P1 (risk)	Fixed - <code>adminMetrics.test.ts</code> , 24 tests
F-07	A table the dashboard reads was never subscribed or published	P2	Fixed - found by the repo's own test during the live run
F-08	11 routes validated the body before checking authentication	P3	Fixed - auth hoisted above validation
F-09	<code>care_plan_reviews.reviewer_id</code> repeats the F-02 contradiction	P3	Fixed - column made nullable
F-10	F-08 was fixed in 11 routes, not in the class; 28 still parsed before authenticating	P3	Fixed - hoisted in all 28, sweep now returns 0
F-11	24 mutating admin routes wrote no audit row at all	P2	Fixed - 23 new actions, one per route, sweep now returns 0
F-12	The test plan referenced five test IDs that were never written, quoted one error string that had changed, and miscounted the admin screens	P2	Fixed - cases written, quotes and counts corrected
F-13	The therapist Overview's primary quick action was a dead anchor	P2	Fixed - points at the route

F-01 - The data reset does not clear four tables, and one of them silently suppresses future test results · P1

Area. Plan §6 (STEP 0), `SETUP-RESET-001`. **Class.** VERIFIED-SOURCE, with schema line numbers.

What the product promises. The reset button's own copy reads: *"Deletes everything - people, sessions, purchases, catalog, settings. Admin logins survive. No undo."* `AGENTS.md` states the rule explicitly: *"Adding a table means adding it to that TRUNCATE list, or a reset silently leaves its rows behind."*

What actually happens. The last definition of `debug_reset_all_data` is at `schema.sql:4876`. These tables were added **after** it and were never added to its `TRUNCATE` list:

Table	Created at	Cleared by the reset?	Why
<code>care_plans</code>	6092	Yes - by CASCADE	References <code>session_entitlements</code> and <code>treatment_categories</code> , both truncated
<code>care_plan_versions</code>	6136	Yes - by CASCADE	References <code>appointments</code> , <code>session_notes</code> , both truncated
<code>contact_reveal_log</code>	6395	Yes - by CASCADE	References <code>appointments</code>
<code>`communication_flags`</code>	6344	NO	Its only FKs are to <code>profiles</code> , and they are <code>on delete set null</code> - so neither the TRUNCATE nor the <code>delete from auth.users</code> reaches it
<code>`risk_rules`</code>	6478	NO	No foreign keys at all
<code>`risk_signals`</code>	6498	NO	Only FK is to <code>risk_rules</code> , which is not truncated. <code>subject_id</code> is a plain <code>uuid</code> with no FK
<code>`risk_reviews`</code>	6539	NO	FKs to <code>profiles</code> and <code>risk_signals</code> , neither truncated

Note that `therapist_schedule_state` **is** in the list despite also being added later, so the list has been maintained for the roster work - the care-plan, evidence and risk tables were simply missed.

Why this is P1 rather than cosmetic. Three consequences, in increasing order of seriousness:

1. **The button lies.** "Deletes everything" is false for four tables.
2. **Orphaned evidence.** `communication_flags` rows survive with `author_id` and `patient_id` set to null by the cascade, so the flag remains but the people it names are gone. The table has a `BEFORE UPDATE OR DELETE` trigger that raises, so **there is no way to clear these rows short of adding the table to the TRUNCATE list** - a row delete is refused by design.
3. **The next test run is silently corrupted.** `risk_signals` carries a partial unique index giving at most one **open or reviewing** signal per `(rule, subject)`. A leftover open signal from a previous run holds that slot, so the same rule firing again on the same subject **writes nothing**. A tester executing `ADM-RISK-001` on a "clean" database sees an empty queue and reports a broken detector. This is exactly the leftover-state trap the plan warns about in §6.1 - except the plan tells the tester the reset protects them from it.

Fix. Add `communication_flags`, `risk_signals` and `risk_reviews` to the `TRUNCATE` list in a new `create or replace` of the function at the end of `schema.sql` (the file's own append-only convention). Decide `risk_rules` deliberately: it is configuration, so either truncate it and let the seed re-insert the eight rules, or reset it to defaults the way `site_settings` is - but do not leave it in its current state, where an edited threshold survives a wipe with nothing telling anyone.

Retest. `SETUP-RESET-001`, plus a new assertion: after a reset, `select count(*) from communication_flags` and `from risk_signals` both return `0`.

FIXED. `debug_reset_all_data` is redefined at the end of `schema.sql` (the file's append-only convention - later definitions win). `communication_flags`, `risk_signals` and `risk_reviews` are now named in the `TRUNCATE`, and `care_plans` / `care_plan_versions` / `contact_reveal_log` are named explicitly too rather than relying on CASCADE, so a future foreign-key change cannot quietly take them back out of the reset. `risk_rules` is deliberately **not** truncated - it is configuration, like `site_settings` - and is instead reset to its seeded defaults, because an empty `risk_rules` would silently disable every detector rather than restoring it. `SETUP-RESET-001` now asserts both counts and the restored thresholds.

F-02 - `risk_reviews.reviewer_id` is NOT NULL with ON DELETE SET NULL · P2

Area. Plan §14.1, `ADM-RISK-002`. **Class.** VERIFIED-SOURCE.

Evidence. `schema.sql:6542`:

```
reviewer_id uuid not null references profiles(id) on delete set null,
```

These two clauses contradict each other. When the referenced profile is deleted, Postgres attempts to set `reviewer_id` to null, which the `NOT NULL` constraint refuses - so **the delete raises** rather than the reference being cleared. A scan of the whole schema found this is the **only** occurrence of the pattern; every other `on delete set null` column is nullable.

Reachability. Reviews are written by admins, and the reset preserves admins, so the standard reset path does not hit it. It becomes reachable when an admin account is deleted at all, and specifically when an admin who has reviewed a risk signal is demoted to another role and a reset then runs: the `delete from auth.users` cascades to `profiles`, the SET NULL fires, the constraint refuses, and **the entire reset transaction aborts** - leaving a half-tested database and an error a tester will not be able to interpret.

Fix. Pick the intent and state it. Either `reviewer_id` is nullable (a review outlives its reviewer, which matches how the other evidence tables treat authorship), or the reference is `ON DELETE RESTRICT` and a reviewer cannot be deleted while reviews exist. The current pair is neither, and produces a failure mode nobody chose.

FIXED. `alter table risk_reviews alter column reviewer_id drop not null`. A review outliving its reviewer is the intended reading - the same posture `communication_flags` already takes with `author_id`, whose reference is nullable for exactly this reason. The reviewer's name is resolved at render time, as it already is for a flag whose author has gone.

F-03 - `pain_assessments` is append-only by RLS alone, unlike its five siblings · P3

Area. Plan §12.4, `THR-HP-005`. **Class.** VERIFIED-SOURCE.

Evidence. `pain_assessments` is protected by `revoke update, delete on pain_assessments from authenticated` (`schema.sql:2212`). Five comparable tables use a **trigger** instead: `session_credit_ledger`, `care_plan_versions`, `communication_flags`, `contact_reveal_log`, `risk_reviews`.

Why the difference matters. The codebase already articulates the reason, in its own comment beside the ledger: *"The revoke covers a browser session; every route in this app writes with the service-role client, which bypasses RLS entirely. For a table whose whole value is that it cannot be rewritten, 'no route updates it' is not the same guarantee as 'an update raises'."* Pain Map data is clinical exam history whose value is

precisely that a re-assessment is a new row, so a trend can be shown against the previous visit. It is protected one grade below the tables that share that property.

Severity is P3, not higher, because no route today updates it, and the gap is a consistency and defence-in-depth issue rather than a live hole.

Fix. Reuse the existing `communication_evidence_is_append_only()` trigger function - it already raises with the table's own name - and attach it to `pain_assessments`.

FIXED. `pain_assessments_no_change`, a `before update or delete` trigger using that same function. The table now carries the same guarantee as its five siblings.

F-04 - The plan overstates the admin scope guard (documentation defect) · P3

Area. Plan §4.4, §15.3, `ADM-SET-027`. **Class.** VERIFIED-SOURCE.

What the plan says. *"Every admin route guards with `requireAdminScope(section)`, not `getAdminUser()`."*

What is true. **92 of 95** admin routes call `requireAdminScope`. Three do not, and all three are deliberate and **stricter**:

Route	Guard	Why stricter
<code>admin/set-admin-scope</code>	<code>getAdminContext</code> + <code>scope !== "full" → 403</code>	A section check would let a <code>finance</code> admin widen its own access
<code>admin/debug-reset</code>	<code>getAdminContext</code> + <code>scope !== "full" → 403</code>	Gate 3 of the reset
<code>admin/create-account</code>	<code>getAdminContext</code> + full-only for minting an admin, full-or-operations otherwise	A scoped admin creating a full admin would hand itself the keys

Impact. No security impact - the exceptions are tighter than the rule. The defect is in the test plan: a tester reading §4.4 literally would file these three as violations. Correct the plan to state the rule as "every admin route guards on scope; three guard on `full` explicitly, because a section check would be too weak."

Fix applied in this session: the wording is corrected in the plan sources.

F-05 - The plan asserted a row cap on All Sessions that no longer exists (documentation defect) · P3 · Fixed

Area. Plan §14.2, `ADM-SESS-001`. **Class.** VERIFIED-SOURCE.

The plan asserted *"at most 200 rows are painted before a 'Show all' affordance appears"*. That was the **previous** design. The screen now uses the shared pager: `usePagedList(rows, { storageKey: "admin-sessions" })` with `DEFAULT_PAGE_SIZE = 10`. `AGENTS.md` records the change explicitly - *"Don't cap a list at an arbitrary number with a 'Show all' escape hatch: that was what All Sessions did, and 'Show all' then painted every row anyway."*

A tester would have reported a working screen as a defect. **Corrected in this session;** the test now asserts the pager, the per-browser page-size memory, and that exports and totals still run over the whole filtered set.

F-06 - The revenue split has no unit test · P1 (risk, not a defect)

Area. Plan §16. **Class.** EXECUTED (by absence) + VERIFIED-SOURCE.

`moneyByBucketFor` in `src/lib/adminMetrics.ts` is the only place the clinic's money is divided, and the plan's two identities live inside it. It feeds the Money summary strip, the tiles and the breakdown chart, so a defect there is wrong on three screens at once and wrong in the same direction, which makes it invisible to cross-checking between them.

There is no `adminMetrics.test.ts`. The nine test files cover the roster, home-visit pricing, the scanner, masking, settings, risk vocabulary, care plans and the consultation rule - every dependency-free module in `src/lib` **except** the one that computes the split. Its own comments record **four** distinct historical misstatements it has already been corrected for: counting uncompleted sessions in the therapist cut, dropping the travel fee into the clinic's share, subtracting refunds from a margin that had already had a cut taken, and collapsing "no hospital" with "hospital share unconfigured". Every one of those is exactly the kind of arithmetic a table-driven unit test pins down permanently.

This is not a defect - the current implementation reads correctly, and the identities hold by construction because the clinic share is computed as the difference rather than accumulated independently. It is the **largest untested risk surface in the application**, and it is the one the finance section of the test plan spends the most manual effort re-deriving by hand.

FIXED. `src/lib/adminMetrics.test.ts` - **24 tests**, structured around the four historical misstatements rather than around the function, so the point is that they stay fixed:

- both identities asserted, including on an empty range;
- a completed paid session earns a cut, a paid-but-undelivered one does not, and a cancelled-and-refunded one earns nothing while still counting in gross;
- a home visit's travel fee lands in the cut and never in gross, falls back to the online share when no home-visit rate is set, and treats a negative fee as zero rather than as a credit;
- the partner's commission is taken on net, so a refund reverses it, while the therapist's is untouched;
- "not referred" and "referred but unconfigured" stay apart - same money, opposite treatment;
- unpaid, slot-less and out-of-range sessions are ignored, and bucketing is by slot time rather than by when the money was taken;
- and the whole seven-row §16.1 reference dataset, so the figures a tester re-derives by hand on the Money screens are the figures pinned here.

The operational rates beside it (`computeNoShowRate` , `computeCancellationRate` , `computeRepeatBookingRate`) are covered too - including that they return **percentages, not fractions**, and `null` rather than zero when there is nothing to divide.

Original recommendation, for the record. Add `adminMetrics.test.ts` covering the §16.1 reference dataset: a completed paid session, a paid-but-not-completed one, a refunded one, a hospital-referred one, one with an unset therapist share, and a home visit with a travel fee - asserting both identities plus the excluded count. That dataset is already written and hand-computed in the plan; it converts directly into a test table. The manual finance tests then become a check on the *screens*, not on the arithmetic.

F-07 - `therapist_schedule_state` was read by the dashboard but never subscribed or published · P2 · Fixed

Area. Plan §14.0 (realtime), `ADM-TODAY-002`. **Class.** EXECUTED - this one was caught by the repository's own test suite during the live run, not by reading code.

What failed. `e2e/admin-multi-admin.spec.ts` H-006 asserts that every base table the admin dashboard queries appears in `ADMIN_REALTIME_TABLES`, and that everything in that list is added to the `supabase_realtime` publication. It failed with one name: `therapist_schedule_state`.

Why the lint check did not catch it. `scripts/check-realtime-coverage.mjs` and H-006 ask different questions. The script checks that every table the UI **subscribes to** is **published** - it passed, because this table was in neither list. H-006 checks that every table the dashboard **reads** is subscribed to, which is the stricter direction. Two checks, one gap between them.

Origin. The table arrived with the roster rebuild (`1aef991`), which added it, queried it from the dashboard at `page.tsx:327`, and updated the reset function's TRUNCATE list for it - but not the realtime arrays. It predates the changes in this branch.

Consequence, stated honestly. Benign but real. The dashboard reads `therapist_schedule_state` to hand the roster editor the version its next save must match. Without a subscription, one admin saving a roster leaves a second admin's open dashboard holding a stale version - the compare-and-swap catches it, so nothing is corrupted, but the second admin gets a 409 telling them to reload where a live refresh would simply have happened. It is a papercut, not a correctness hole; it is a finding because the codebase's own stated rule is that a new table goes into one of the two arrays **and** gets its `alter publication` line in the same change.

FIXED. Added to `ADMIN_REALTIME_TABLES` (the operational channel - a roster change is operational) and given the guarded `alter publication` block every other publication line in `schema.sql` uses. H-006 re-run: **passes.** `check:realtime` now reports 39 subscribed tables, all published.

F-08 - Eleven routes validated the request body before checking who was asking · P3 · Fixed

Area. Plan §18.2, `SEC-ROUTE-002`. **Class.** EXECUTED.

What was found. Calling 25 routes with no cookie and an empty body, 11 answered **400 with a field-validation message** rather than 401/403:

```
/api/appointments/cancel      400 {"error":"Missing appointmentId"}
/api/patient/condition-profile/submit 400 {"error":"Missing data"}
/api/razorpay/create-order    400 {"error":"Missing appointmentId"}
...and eight more
```

How serious, established rather than assumed. The obvious question is whether authentication was enforced *at all*, so the same 11 were re-called with well-formed bodies. Nine then returned 401/403. The two home-visit order routes validated more deeply first - `{"error":"A street address is required."}` - and only refused once a complete address was supplied. So **authentication was always enforced; nothing could be done anonymously.** No data leaked, no action was possible, and no response contained internals.

That makes this an **ordering** defect, not a hole - which is why it is P3 and not higher. It is still worth fixing on two grounds: an unauthenticated caller should not drive a route's parsing at all, and the other ninety-odd routes in this application already check auth first, so these eleven were the inconsistent ones.

FIXED. The `createClient` / `getUser` / `if (!user)` block was hoisted above body validation in all eleven, with a comment at each explaining why the order matters. `npm run verify` is green, and the five spec files covering those routes were re-run afterwards: **61/61 passed**, no regression. Re-running the probe: **11/11 refuse an anonymous caller**, and all 25 routes now refuse with nothing leaked.

F-09 - `care_plan_reviews.reviewer_id` repeats the contradiction F-02 removed · P3 · Fixed

Area. Plan §14.2, `ADM-CARE-004`. **Class.** VERIFIED-SOURCE.

Evidence. `schema.sql:7309`, in a table written **after** F-02 was fixed:

```
reviewer_id uuid not null references profiles(id) on delete set null,
```

Pass 1 said of the identical line in `risk_reviews`: *"A scan of the whole schema found this is the only occurrence of the pattern."* It was, on 2 September. The care-plan review step then shipped, its table was written from the older one, and the pattern came back - which is the ordinary way a fixed defect returns: not by being undone, but by being copied from something that predates the fix.

Why it matters. The two clauses cannot both hold. Deleting the referenced profile makes Postgres attempt a null write the `NOT NULL` refuses, so the **delete raises** and the reference is never cleared. The failure surfaces far from here - as an unrelated account deletion aborting with a message naming a table the person deleting has never heard of.

Reachability is narrower than F-02's, which is why this is P3 rather than P2: `care_plan_reviews` is in the reset's `TRUNCATE` list and is emptied before `delete from auth.users` runs, so the standard reset path cannot hit it. It needs an admin profile deleted outside a reset - and this product suspends rather than deletes, deliberately, so that is a hand-run statement rather than a screen.

Fix. `alter table care_plan_reviews alter column reviewer_id drop not null`, appended at the end of `schema.sql` in the file's own guarded, append-only style, with the reasoning in a comment above it. Nullable is the intended reading, exactly as for `risk_reviews` and `communication_flags.author_id`: a decision outlives the person who made it, the reviewer's name is resolved at render time, and the row itself is still unrewritable because `care_plan_reviews_no_change` refuses every update and delete.

Retest. `ADM-CARE-004` is unchanged in behaviour. The storage-layer check is a scratch-database `alter`, then `delete from profiles where id = '<a reviewer>'` - it must not raise.

F-10 - F-08 was fixed in the eleven routes a probe reached, not in the class · P3 · Fixed

Area. Plan §18.2, `SEC-ROUTE-002`. **Class.** VERIFIED-SOURCE (pass 1's evidence was EXECUTED; no server was available this pass).

What was found. Pass 1 called 25 routes with no cookie, found 11 that answered a body-validation `400` instead of a refusal, and hoisted the auth check above validation in those eleven. A static sweep of **all 160** route handlers this pass - comparing the line at which the body is first parsed against the line of the first authentication call - found **28** routes still parsing first, of which **23** return a `400` before anyone has been identified. Among them: `therapist/set-on-leave`, `therapist/session-notes/submit`, `therapist/pain-assessments/submit`, `patient/medical-documents/delete`, `home-visit/verify` and `care-plan/verify`.

Severity is unchanged from F-08, and for the same reason. Authentication was still enforced before anything was written in every one of the 28; nothing could be done anonymously and nothing leaked beyond the shape of the request. It is an ordering defect. What makes it worth its own entry is not the severity but the **shape of the pass-1 fix**: the eleven were the instances the probe happened to reach, and fixing instances rather than the class leaves a rule that looks enforced - the sweep, not the probe, is what says whether it is.

Fix. The `createClient` / `getUser` / `if (!user)` block hoisted to the first statement of the handler in all 28, each carrying the same comment pass 1 wrote for the first eleven. `npx tsc --noEmit` , `npm run lint` , 384 unit tests and a full production build are green on the result, and the sweep now returns **0 of 160**.

Retest. `SEC-ROUTE-002` , extended: call **every** POST handler with no cookie and an empty body, and assert none answers `400` . That sweep is cheap enough to be the regression test this finding needs.

F-11 - A quarter of the mutating admin routes wrote no audit row, including the one that changes a patient's sign-in email · P2 · Fixed

Area. Plan §15.3, `ADM-SET-033` . **Class.** VERIFIED-SOURCE.

What the codebase promises. `AGENTS.md` states it without qualification: *"Every mutating admin route records what happened via `recordAdminActivity()` ... Adding an action without a caller, or a mutating route without a call, puts the log back where it was."* Pass 1 checked one half of that rule - every declared action has a caller - and found it clean. It did not check the other half.

What was found. Of the 102 admin route handlers, three are reads (`export-pdf` , `package-purchase-detail` , `home-visit-purchase-detail`). Of the 99 that write, **24 recorded nothing at all**:

Route	What it changes with no trace
<code>update-patient-contact</code> , <code>update-therapist-contact</code>	The email that gates sign-in - the route updates <code>auth.users</code> as well as <code>profiles</code>
<code>update-patient-notes</code> , <code>update-therapist-notes</code> , <code>update-therapist-display-content</code>	Internal notes and public write-ups
<code>set-therapist-team-visibility</code> , <code>set-therapist-rating-visibility</code> , <code>set-ratings-visible-publicly</code>	Whether a clinician and the clinic's ratings appear on the public site
<code>clear-session-rating</code> , <code>exclude-session-rating</code>	A rating, and therefore the public average
<code>update-home-visit-address</code>	Where a therapist is sent
<code>condition-access/decide</code> , <code>condition-requests/decide</code> , <code>condition-requests/direct-edit</code>	Who may read a patient's record, and what it says
<code>pain-assessments/submit</code>	A clinical exam, written by an admin
<code>intake-questions/update</code> , <code>pain-map-questions/update</code>	The wording of the clinical instruments themselves
<code>review-risk-signal</code>	A decision about a colleague
<code>open-meet-access</code> , <code>retry-meet-sync</code>	A session's meeting
<code>update-lead-status</code> , <code>update-home-visit-waitlist-status</code> , <code>set-referral-capacity-note</code>	Queue state
<code>debug-reset</code>	The whole database

Several have a private trail of their own - `condition_change_requests` records its reviewer, `risk_reviews` is append-only - but none of them reaches the one log an admin can read across every screen, which is exactly what "who changed this?" is asked of.

Why P2. No money moves through any of these, which is why it is not higher. But the audit log exists because *"a refund, a payout settlement or a cash-remittance confirmation was unattributable the moment more than one person had admin access"* - and by that reasoning, a patient's sign-in email being changed

by one of four admins with no record is the same failure in a place that decides whether somebody can get into their own account.

Fix. 23 new actions on `AdminActivityAction`, one `recordAdminActivity` call per route, placed after the write and after any compare-and-swap, best-effort as the rule requires. `ADMIN_ACTIVITY_LABELS` is a `Record<AdminActivityAction, string>`, so the compiler refused the change until every new action had a human label - the mechanism the file's own comment describes, working. Two placements are deliberate rather than mechanical: `retry-meet-sync` logs only on the branch that actually created an event, not on the no-op return for an appointment that already had a link; and `debug-reset` logs **after** the wipe, because the wipe truncates `admin_activity_log` and a row written before it would be erased by the action it describes. The sweep now returns **0 of 99**.

Retest. `ADM-SET-033`, which gained a step 2a naming ten of the newly-logged actions, and an assertion that a freshly reset database's log holds exactly one row: `Reset all data`.

F-12 - The plan referenced five test cases that were never written · P2 · Fixed

Area. The plan itself (§3, §7.4, §24). **Class.** VERIFIED-SOURCE.

What was found. Cross-referencing every test ID the plan *cites* against every ID it *defines*: `PAT-DASH-001`, `PAT-DASH-002`, `THR-DASH-001` and `THR-PROF-001` are named as the coverage for four routes in the §3 route map and are defined nowhere, and `PAT-SESS-005` is named in the simulated-clock scenario table. A route map whose rightmost column names a test that does not exist is worse than a blank cell: **it reads as covered**, and the four routes in question are the landing screens every signed-in session starts on.

Three smaller documentation defects came out of the same pass:

- The plan quoted `complete-session`'s refusal as *"You can mark this done once the session's join window has opened."* The route says *"This session hasn't started yet. You can mark it done once it's under way."* A tester matching the quoted string would file a defect against correct behaviour.
- The Settings split and the User Access rename took the admin dashboard to **31** screens. The §3.6 table lists all 31; its heading and both coverage lines in §24 still said 28, and the Offers & Discounts row pointed at `ADM-SET-016`, which is the sign-out banner.
- `SEC-API-*` was cited as a family of API-level security tests. No such family exists; those calls live in `SEC-ROUTE-002`, `SEC-ADMIN-002` and `SEC-TAMPER-*`.

Fix. The four missing cases written click-by-click - the patient Overview and booking hub, the therapist Overview and Edit Profile - `PAT-SESS-005` repointed at `PAT-SESS-003`, the error string corrected, the counts and the mapping fixed, and `ADM-PEOP-003` rewritten at step level while its screen was open. The plan now defines **358** cases and every ID it references resolves.

Retest. The cross-reference itself: extract every `XXX-YYY-NNN` the sources cite and every one they define; the difference must be empty but for the four documented `SETUP-*` aliases.

F-13 - The therapist Overview's primary quick action was a dead link · P2 · Fixed

Area. Plan §12.1a, `THR-DASH-001`. **Class.** VERIFIED-SOURCE.

What was found. Writing the missing `THR-DASH-001` found what it exists to catch. The therapist Overview's first and only `primary` quick action, **Set your availability**, pointed at `/therapist/dashboard#availability` - an anchor from when the dashboard was one long scrolled page.

Availability became its own route when the dashboards were split, and `dashboardNavItems.ts` has pointed at `/therapist/dashboard/availability` ever since. The action reloaded Overview and did nothing.

Why P2 for a one-line fix. It is the primary call to action on the screen a therapist lands on, for the one task the clinic most needs them to do - an empty roster is why sessions sit unassigned. It also fails silently: the page reloads, so it reads as a slow app rather than a broken link, which is why it survived. `AGENTS.md` names this exact failure mode for admin `?section=&tab=` links (*"a stale link looks like it works - it just quietly lands somewhere else"*); the same hazard applies to any hand-written dashboard href, and this is the only one left in the four dashboards.

Fix. The href points at the route, with a comment saying why the anchor form is wrong here. A sweep for `href: "</role>/dashboard#..."` across all four dashboards returns nothing else.

Retest. `THR-DASH-001` step 7: the tap must land on `/therapist/dashboard/availability` with the weekly schedule editor on screen.

5. What only a live run can establish

The following are **NOT-VERIFIABLE** here and carry real residual risk. They are the areas to run first when an environment exists.

Area	Why source inspection is insufficient
The whole payment funnel end to end	Signature verification, checkout rendering, callback timing and the webhook race are runtime behaviours. The mechanism is right; whether it fires correctly is untested
Concurrency (<code>PAY-CONC-001</code> , <code>THR-AVAIL-004</code> , <code>FIN-PAY-002</code>)	Row locks and CAS predicates are present in the SQL. Whether they actually serialise 12 concurrent reserves needs 12 concurrent reserves
Google Calendar / Meet	Neither the success path nor the capped-retry path can be exercised without credentials
Every screen rendering at all	A build proves it compiles, not that it paints. No page was loaded in this audit
Mobile at 390 × 844, keyboard nav, focus traps, reduced motion	Entirely runtime
Copy fidelity	The plan quotes route-handler strings verbatim; where a component re-words one before display, only a browser will show which the user sees
RLS behaving as written	Policies exist in the file. Whether the live project has them applied is a different question - and the schema-apply workflow is the only thing that closes it
The reset actually running	F-01 was found by reading the function. It should be confirmed by running it and counting rows
The audit rows F-11 adds	The 23 new <code>recordAdminActivity</code> calls compile and are placed after their writes. Whether each one lands a readable row is <code>ADM-SET-033</code> step 2a, and needs a database
The 28 hoisted routes still behaving	Type-checked, linted, built and unit-tested. The five spec files covering them were not re-run this pass - there was no server to run them against - so <code>SEC-ROUTE-002</code> and the therapist/patient specs are the first thing to run when one exists

6. Verdict and recommended sequence

Pass, after remediation, on both passes. The invariants the plan cares most about - the capture path, the ledger constraints, the append-only triggers, the split identities, the scope guards, the whole-hour rule, the single checkout quote - are all genuinely implemented where the plan says they are. **All thirteen findings are fixed**, including the three documentation defects in the plan itself, and `npm run verify` is green: 384 unit tests in 23 files, lint, and a full production build.

What pass 2 changes about how to read this report. Two of its five findings are the same shape: a rule that *was* fixed, in the instances somebody looked at. F-08's ordering fix reached the eleven routes a probe called, and 28 others kept the old order; F-02's schema fix was correct and a table written afterwards copied the pre-fix pattern. Neither was a regression in the ordinary sense - nothing was undone - and neither would have been caught by re-running the check that found it the first time. **The sweeps in §3.13 are the answer to that**, and they are cheap: each one is a few lines of static analysis over the whole repository, and each is worth re-running before every release rather than once per audit.

One thing still needs a person, not a commit. The schema changes across both passes - the redefined reset function, the two nullable `reviewer_id` columns, the `pain_assessments` trigger, the `auto_assign_therapist_enabled` column and everything the branch has added since - reach a live database only when `supabase/schema.sql` is applied, either by hand with `node scripts/run-schema.mjs` or by the schema-apply workflow on a push to `main`. **Until that runs, the fixes exist in the file and not in the database**, which is the exact failure mode the schema conventions warn about: the app looks fixed in review and still behaves the old way in production. Apply it, then re-run `SETUP-RESET-001` and confirm the two counts come back zero.

Then run, in this order: the §16.3 payment-integrity sweep (highest value per hour, and entirely unverifiable statically), the §21 cross-role checks (they catch disagreements no single-role test can), then the §18 security sweep at the route level, then everything else in the plan's own recommended order.

7. Product review - flows worth changing

This section is written from a product rather than a QA seat. Each item names what the code does today, what it costs, and what changing it would buy. They are ordered by expected value, not by effort.

Seven of the nine shipped in the same pass as the findings. Each carries its outcome at the end. The two that did not are 7.6 and half of 7.7: both are owner decisions about *policy* rather than defects, and one of them (the ledger flip) is explicitly conditioned on evidence that does not exist yet.

#	Item	Outcome
7.1	Auto-assign a therapist at payment	Shipped , behind a switch, off for one release
7.2	Name who is waiting for a recommendation	Shipped
7.3	Payment reassurance on the first failure	Shipped
7.4	Webhook secret on System Health	Shipped
7.5	Sweeps only run when an admin looks	Documented - an ops change, not a code one

#	Item	Outcome
7. 6	Patient approval queue	Half shipped - the screen now says what it decides; removing the gate is the owner's call
7. 7	Two dark features	Suggestions on by default. Ledger deliberately left off
7. 8	Smaller items	Shipped (page size, waiting-state date); therapist timezone still open

7.1 The funnel stalls on a manual admin step, and the data to remove it already exists · High value

Today. A patient picks a time, pays, and the appointment sits `requested` with `therapist_id = null` until an admin opens the dashboard and assigns someone. Only then is it `confirmed`, only then does a Meet link exist, and only then does the therapist see it. The roster - weekly schedule, exceptions, leave - is maintained in detail and deliberately does **not** feed the patient's picker.

What it costs. The gap between "patient has paid" and "patient has a confirmed session with a named clinician" is bounded by how quickly a human opens a browser. At a weekend or overnight booking that is hours. The patient's own screen says `Requested` with no therapist, which is the least reassuring possible state immediately after paying.

The change. Keep the picker unfiltered - that separation is deliberate and correct, and constraining what a patient may ask for is a different, worse product. But use the roster **server-side at payment**

confirmation: if exactly one approved, active, non-on-leave therapist is free for that slot per the existing conflict check, assign them automatically and confirm. If zero or more than one, fall through to the queue exactly as today.

Why it is safe. Every component already exists - `therapistAvailability.ts`, `checkTherapistConflict.ts`, and the auto-assign path used by `bookPackageSession()`, which already assigns, confirms and mints a Meet link without an admin. This is wiring two existing mechanisms together, not new logic. The admin keeps full override.

Measure. Median minutes from `paid_at` to `status = 'confirmed'`, before and after.

SHIPPED. `src/lib/autoAssignTherapist.ts`, called from `/api/razorpay/verify` and `/api/razorpay/webhook` - both, so a patient who pays and closes the tab gets the same outcome as one who waits for the page, and neither path can develop its own idea of who is free. It reads the roster (weekly template, that date's exceptions, `on_leave`) plus the existing conflict check, and assigns only when **exactly one** eligible therapist is free, or when the patient's own `preferred_therapist_id` is among the free ones. Zero or two-or-more does nothing and the session waits in the queue exactly as before. It never throws. `decideAutoAssignment()` extracts that rule with the database taken out, so the judgement itself is unit-tested (10 tests) rather than only reachable through a payment. Gated by `auto_assign_therapist_enabled` - **off for its first release**, because it changes what happens to a booking after money has moved and the honest way to introduce that is with the previous behaviour one click away. `ADM-SET-021` covers all twelve cases.

7.2 Consultation → programme conversion depends on a therapist remembering · High value

Today. A programme can only be bought from a care plan, and a care plan can only be written from the session-note dialog of a **completed** session the therapist ran. The nudge is passive: a "Notes to write"

figure on the therapist's Overview and a feed item. The detector that would measure the failure, `plan_conversion_low`, **ships disabled**.

What it costs. The entire revenue model above a single consultation runs through one optional action by a busy clinician at the end of a session. Nobody currently knows the conversion rate, because the rule that would compute it is off.

The change, in two steps and in this order. First **measure**: enable `plan_conversion_low` in a log-only posture for a month to establish the clinic's baseline - the reason it ships disabled is precisely that a threshold invented before a baseline fires on everyone or nobody, and that reasoning is right. Then, and only then, decide whether the prompt needs strengthening.

A cheap intervention that does not need the baseline: the session-note dialog already contains the recommendation panel. Make "no recommendation written" an explicit `needsYou` feed item with the patient's name and the date of the session, rather than an aggregate count. A named item is acted on; a number is not.

SHIPPED (the cheap half). The therapist's Overview feed now carries a named item per patient - "QA Patient A is waiting to hear what next" - pinned by `needsYou` and linking to that patient's chart. A patient with a live or already-purchased recommendation is excluded; a declined or withdrawn one is included, because that thread is open again. Capped at four, most recently seen first, beside the existing note nudge. `THR-CARE-006` covers it. **The measurement half is still the right first move** - `plan_conversion_low` remains disabled until the clinic has a baseline, and that reasoning has not changed.

7.3 The payment reassurance arrives two failures too late · Cheap, do it now

Today. After a failed or dismissed payment the wizard shows an error and re-labels the button **Pay ₹...**
Now. The reassuring message - "Your booking is saved as pending - you can come back and pay any time from your dashboard" - appears only after **three** failed attempts.

What it costs. A patient whose card fails once has no idea their booking survived. The most likely next action is to close the tab, and they have no reason to believe anything is waiting for them.

The change. Show the "your booking is saved" line on the **first** failure. Keep the escape-hatch link at three. One conditional; no new state.

SHIPPED. One conditional in `BookingWizard.tsx`: from the first failure the card reads "Nothing was lost - your booking is saved and still held as unpaid. You can try again above, or pay later from your dashboard." The amber escape hatch to the dashboard still waits until three, because that is the point at which retrying here has plainly stopped working. `PAT-PAY-003`.

7.4 Losing money silently when one environment variable is unset · Cheap, high consequence

Today. Without `RAZORPAY_WEBHOOK_SECRET`, `/api/razorpay/webhook` answers `503`. A patient who pays and closes the tab before the browser callback lands leaves a **paid Razorpay order against an unpaid booking**. Nothing anywhere reports this.

What it costs. Money collected with no session against it, discovered only when a patient complains or someone reconciles Razorpay by hand.

The change. Settings → System Health already exists and already surfaces sync failures and accounting disagreements. Add one row: webhook secret configured, yes or no, red when no. It is the only configuration whose absence silently loses money, and the screen that should say so is already built.

SHIPPED. A **Payment Confirmations** panel at the top of System Health. Configured, it states that a payment is confirmed by whichever arrives first, browser or webhook. Unconfigured, the whole panel turns

red and says plainly that a patient who pays and closes the tab will leave a paid order against an unpaid booking with nothing flagging it, and names the variable to set. The presence of the secret is read server-side in the dashboard page and only the boolean crosses to the browser. `ADM-SET-030`.

7.5 Time-based work only happens when an admin is looking · Medium

Today. There is no cron, by design. The sweeps run at page render - and the audit confirms **the failed-Meet-sync retry and the risk detector run only on the admin dashboard render**. Package expiry additionally runs on the patient dashboard.

What it costs. On a quiet day where no admin opens the back office, no failed Meet link is retried and no risk signal is detected. The system's self-healing is coupled to somebody's browsing habits.

The change. Do not build a worker - the no-cron rule buys real simplicity and the sweeps are correctly bounded. Instead point a free uptime monitor at one authenticated-free endpoint every fifteen minutes. That is a configuration change with no code, and it converts "when an admin looks" into "every fifteen minutes" for the two sweeps that most need it.

NOT SHIPPED, deliberately - this one is an ops change, not a code change. Adding an endpoint that runs the sweeps would be a worker in all but name, and would undo the property that makes the no-cron design defensible: every sweep is bounded because it runs inside a render somebody is waiting for. The recommendation stands as written, for whoever configures the deployment.

7.6 Patient approval may be a queue that no longer earns its keep · Worth a decision

Today. Patients are gated on `approved`. But a genuine payment attempt auto-approves them. So the patient approval queue only ever contains people who registered **without** paying - and the plan's own \$4.3 explains that this is deliberate, so that a failing card does not bounce someone to `/pending-approval`.

The question for the owner. What is the queue actually filtering? Everyone who pays is approved automatically; everyone who does not pay cannot book. The remaining population is people who registered to browse. If the answer is "nothing much", the queue is manual work with no decision attached to it, and it delays the one group who took a deliberate action but have not yet paid.

Two honest options. Keep it and state its purpose in one sentence on the screen so a new admin knows what they are deciding; or drop patient approval to "review only if flagged" and keep the human gate for therapists, where credential checking is a real decision.

HALF SHIPPED. The first option is done: Today → Approvals now states what it is deciding - *"Therapists here are waiting on a credentials check. Patients here registered without booking - a patient who starts a payment is approved automatically, so approving one from this list only affects what they can see, never whether they can pay."* **The second is not, and should not be done by an engineer.** Removing a gate on who can use a clinical product is a policy decision with a compliance dimension, and the sentence above is what makes it possible to take that decision with the facts in view.

7.7 Two features are built, dark, and rotting · Worth a decision

`therapist_suggestions_enabled` defaults **false**. `entitlement_ledger_authoritative` defaults **false**.

The first is the main re-booking mechanism - a therapist proposing the next session on a live programme - fully built, tested, and switched off. The second is a completed migration to an append-only credit ledger, with both systems written in parallel and a verification report on Settings → System Health.

Both defaults were right when written. Both now need a date rather than a default:

- **Suggestions:** launch it or delete it. A dark feature accumulates maintenance cost and drifts out of test coverage.
- **Ledger:** the parallel-write period is the risky state, not the destination - two sources of truth for session balances is precisely the condition the ledger was built to end. Once System Health has been clean over real traffic for an agreed window, flip it, then plan the removal of the counter writes. Carrying both indefinitely is the one outcome nobody chose.

7.8 Smaller items

Item	Today	Suggested
All Sessions page size	Defaults to 10 rows	Shipped at 25. The per-browser <code>storageKey</code> means anyone who already chose a size keeps it
A patient locked out of their own health profile	Read-only until a therapist writes the first record, with the CTA absent rather than disabled - correct, and deliberately quiet	Shipped. The waiting panel now names the session by date - <i>"Your session on 11/09/2026 is when this gets filled in"</i> , or after it <i>"Expected at your session on ... If it still isn't here in a day or two, tell us and we'll chase it."</i> A patient with no session yet sees no date line rather than a placeholder
Therapist timezone	<code>profiles.timezone</code> is displayed on the roster and settable nowhere	Still open - it is a product decision (is this a one-timezone clinic?), not a defect, and is item 2 in the plan's own clarifications list. Either expose it on the therapist's profile or drop the column
`risk_rules` after a reset	Survives with edited thresholds (F-01)	Whatever is decided for F-01, make it deliberate - configuration that survives a wipe should do so for the same stated reason <code>site_settings</code> does

7.9 What is working well, and should not be traded away

A product review that only lists changes misrepresents the system. These are load-bearing decisions that are correct and should survive future pressure to "simplify":

- **One capture path.** `record_payment_capture` being a single idempotent database function is why duplicate webhooks, gateway retries, a webhook racing a callback and a double-clicked Pay button are all safe without any of them knowing about each other. Do not add a second fulfilment path; extend its `purpose` check.
- **A therapist picks a package, never a price.** "The therapist set their own price" is not a policy someone enforces - it is a thing the schema cannot express. That is a much stronger guarantee than a rule, and it costs nothing.
- **Evidence tables append-only by trigger, not RLS.** Every route writes with the service role, which bypasses RLS entirely; a record the evidenced party could edit is not evidence. The distinction is subtle and correct.
- **A flag is never an accusation.** The risk queue deliberately carries no action buttons, and stores the row ids behind a finding rather than a score. That is what makes running heuristics over clinical data defensible at all.
- **The roster does not filter the patient's picker.** It reads like an omission and is a decision. Anyone proposing to "fix" it is proposing a product change with a deploy-sized blast radius.

7.10 Where the pass-1 recommendations stand (checked 4 September 2026)

Item	Then	Now
7.1 The funnel stalls on a manual admin step	Every paid session waited for an admin to assign a therapist	Shipped. <code>autoAssignTherapist.ts</code> assigns when exactly one rostered, approved, un-clashing therapist is free - or when the patient's own requested therapist is among them - from both payment-confirmation paths, gated by <code>auto_assign_therapist_enabled</code> and failing closed. Still off by default , which is right for its first release and is now the thing to put a date on
7.2 Conversion depends on a therapist remembering	Nothing prompted a recommendation after a session	Partly shipped. The therapist's Overview counts notes owed and the feed carries a <code>needsYou</code> item; the clinic's review queue ages in words and colours anything past four hours. There is still no prompt that says "this patient finished a consultation and has no recommendation"
7.3 Payment reassurance two failures late	Reassurance at three failures	Shipped , at the first
7.4 A missing webhook secret loses money silently	Nothing reported it	Shipped - the Payment Confirmations panel on System Health
7.5 Time-based work happens only when an admin looks	Lazy sweeps at render	Unchanged, and still correct for this deployment. Worth revisiting only if a scheduler ever exists
7.6 Patient approval may no longer earn its keep	A queue in front of every new patient	Still open - a product decision
7.7 Two features built and dark	Suggestions and the ledger both defaulted off	Half resolved. <code>therapist_suggestions_enabled</code> now defaults true on a fresh database (an established clinic keeps its stored value until an admin toggles it). <code>entitlement_ledger_authoritative</code> is still false , and the parallel-write period is still the risky state rather than the destination - this is the one item on the list that has been waiting longest and costs the most to keep waiting

The one new item this pass adds is not a feature: it is that two of the five findings were fixed-then-recurred (F-09, F-10), and the cheapest defence is the sweep set in §3.13 run before a release rather than a rule written in a document. A rule catches the people who read it; a sweep catches the code.